

Theorem 14.2

The language L is recursive if and only if L and \bar{L} are recursively enumerable.

Proof

Let L be any recursive language. Then, by definition, L has a Turing machine for it that halts on all inputs. Since L has a Turing machine for it, L is also recursively enumerable. By Theorem 14.1 we know that \bar{L} is also recursive and thus also recursively enumerable. We have thus shown that if L is recursive, then L and \bar{L} are both recursively enumerable.

Now we will assume that L and \bar{L} are both recursively enumerable. Then, by definition, L has a Turing machine that accepts it, call it M , and \bar{L} has a Turing machine that accepts it, call it \bar{M} . We can design a Turing machine M' for L that halts on all inputs as follows. On any input, M' will simulate M on the input and in parallel simulate \bar{M} on another copy of the input (perhaps on a second tape), alternating steps for these two simulations. We know that we can have a Turing machine simulate another Turing machine based on the existence of the Universal Turing Machine. If the simulation of M ever halts and accepts, then M' will halt and accept. If the simulation of \bar{M} ever halts and accepts, then M' will reject. Notice that M' accepts exactly the language L since M is a Turing machine for L and \bar{M} is a Turing machine for \bar{L} . Moreover, exactly one of M and \bar{M} must halt and accept because the input is either in L (in which case M will accept) or it is not (in which case \bar{M} will accept). Since M' is a Turing machine that accepts L and halts on all inputs, this means that L must be recursive.

We have thus proved that L is recursive if and only if L and \bar{L} are recursively enumerable.

Notice in the above proof that we had to run the simulations for M and \bar{M} in parallel because it could be possible for either M or \bar{M} to run forever on rejection and so if we had run them sequentially the computation may never end.

14.6 A NON-COMPUTABLE PROBLEM

To show that there are non-recursive problems, we first need to understand that all strings and Turing machines can be enumerated (listed in a numbered list). To list off all the strings in a fixed order, we can use shortlex ordering. For a review of shortlex ordering of strings, take a look at Section 1.10. To list all Turing machines, we can write an encoding for each one (see Section 14.4) and then also use a shortlex ordering on these. Let us call the Turing machines generated in this manner M_i for each string i .

We are now ready to define the Diagonal Problem (D). The Diagonal Problem can be intuitively thought of as the set of strings that when interpreted as a Turing machine accepts itself as an input.

Definition 14.4

$$D = \{i : M_i \text{ accepts input } i\}$$

To understand the above definition, we have to interpret the string i in two distinct ways. In one interpretation, we treat the string i as the encoding for a Turing machine M_i as we saw in Section 14.4. If the string i does not correspond to the encoding of any Turing machine, then we simply define M_i to be a simple Turing machine that rejects all inputs by immediately halting and rejecting. The second interpretation of i is simply that of a string—an input string to be accepted or rejected.

Next, we define the complement of the language D , called \overline{D} .

Definition 14.5

$$\overline{D} = \{i : M_i \text{ rejects input } i\}$$

We will now show that \overline{D} does not have a Turing machine that accepts it. Before formally proving this, let us build some intuition about why this is the case. We use a technique called **diagonalization**.

Imagine that we could construct an infinitely large table that has all the Turing machines on the rows and the strings on the columns, as illustrated in Figure 14.3. For simplicity, we will assume that the strings are over the alphabet $\{0, 1\}$. Each entry in the table will indicate whether the Turing machine on the row accepts (A) the string on the column or rejects (R) it. (The entries in Figure 14.3 are not real and just shown for illustrative purposes.) We will focus on the diagonal entries of the table (marked with boxes) as the diagonal language and its complement are about these values. We can show that M_λ is not a Turing machine for \overline{D} because (in this example table) λ is accepted by M_λ , putting λ in D (by the definition of D) but then a Turing machine for \overline{D} should reject λ . A similar argument shows that M_0 is also not a valid Turing machine for \overline{D} . On the other hand, the Turing machine M_1 rejects the string 1, which means that 1 is in \overline{D} , but now M_1 cannot be a Turing machine for \overline{D} as it should have accepted 1. Continuing this argument for each diagonal entry (whether it is A or R), shows that no Turing machine in our list of all possible Turing machines can accept \overline{D} . This shows that \overline{D} is not recursively enumerable.

	λ	0	1	00	01	...
M_λ	A	R	R	R	R	...
M_0	A	A	A	A	A	...
M_1	R	R	R	R	R	...
M_{00}	R	A	A	R	A	...
M_{01}	A	A	R	R	A	...
...

Figure 14.3: Table of whether each Turing machine accepts/rejects each string

Theorem 14.3

The language \overline{D} is not recursively enumerable.

Proof

We will prove that there cannot exist any Turing machine that accepts \overline{D} via contradiction. Assume that \overline{D} is accepted by a Turing machine with encoding i . We have two possible cases:

M_i accepts i : If it is the case that M_i accepts i , then by the definition of D it must be that $i \in D$. Then M_i cannot be a Turing machine for \overline{D} since it should reject strings in D such as i .

M_i rejects i : If it is the case that M_i rejects i , then by the definition of \overline{D} it must be that $i \in \overline{D}$. Then M_i cannot be a Turing machine for \overline{D} since it should accept strings in \overline{D} such as i .

We get a contradiction in both cases. Thus, a Turing machine for \overline{D} cannot exist and so \overline{D} is not recursively enumerable.

We can conclude from this the following result for D .

Theorem 14.4

The language D is not recursive.

Proof

We showed earlier in Theorem 14.2 that a language is recursive if and only if the language and its complement are recursively enumerable. Since we know that the complement of D is not recursively enumerable, it must be that D is not recursive.

We have shown that \overline{D} and D are not recursive or computable, but they are not very natural languages. We will use them to next show that other, more easy to understand problems, are not computable as well.

14.7 REDUCTIONS

We will now show how to prove the non-computability of another language by bootstrapping from the previous two languages using a technique called a reduction. Consider the following language:

Definition 14.6

$A = \{(i, x) : M_i \text{ accepts } x\}$.

The language A consists of all pairs of Turing machine encodings and strings such that the Turing machine accepts the string. We encode pairings such as this by

introducing new symbols $(,)$, and $,.$ If we could compute this language, we could then tell whether any arbitrary program accepts an arbitrary input. Unfortunately, no Turing machine (that halts on all inputs) exists for it as we see next.

Theorem 14.5

The language A is not recursive.

Proof

We will assume that A is recursive for a contradiction. Then there exists a Turing machine M for A that halts on all inputs. We will use M to design a Turing machine for D that halts on all inputs.

We design a Turing machine M' that on input i will run M on the input (i, i) and then do what it does (accept if M accepts or reject if M rejects). We know that it is possible to design a Turing machine that runs another machine based on what we learned about the Universal Turing Machine.

We now show that M' is a Turing machine for D that halts on all inputs. This has three parts:

The Turing machine M' must halt on all inputs as the machine M halts on all inputs.

If $i \in D$, then M_i accepts i (by the definition of D), so M accepts (i, i) and thus M' accepts.

If $i \notin D$, then M_i rejects i (by the definition of D), so M rejects (i, i) and thus M' rejects.

The last two parts together give us that $i \in D$ if and only if M' accepts i , which means that M' correctly accepts all strings in D and rejects all others, making it a Turing machine for D that halts on all inputs. However, we already proved that D is not recursive, which gives us a contradiction. Thus, A is not recursive.

Intuitively, the above reduction uses the fact that determining whether M_i accepts an arbitrary string x should be even harder than determining whether M_i accepts the specific string i .

Since we believe Turing machines to be able to perform any mechanical computation possible on a computer, we can rephrase the implications of the above theorem thus:

It is impossible to design a program that can predict the output of an arbitrary program on arbitrary input.

14.8 PROGRAM COMPARISON

We'll next see a reduction that is a little more involved. The notation $L(M_a)$ denotes the language of the Turing machine M_a .

Theorem 14.6

The language

$$SAME = \{(a, b) : L(M_a) = L(M_b)\}$$

is not recursive.

The language *SAME* is conceptually the pairs of Turing machines that behave identically on all inputs. Translated from Turing machines to more practical usage, the uncomputability of *SAME* implies the following:

It is impossible to design a program that can tell if two programs behave identically on all inputs.

By showing that this problem is non-computable, this means that it is impossible for your CS instructor to design a program that takes as input two programs (for example, a working solution and a student's submission) and determine whether they behave identically on all inputs.

Proof

We will assume that *SAME* is recursive for a contradiction. Then there exists a Turing machine *M* for *SAME* that halts on all inputs. We will use *M* to design a Turing machine *M'* for *D* that halts on all inputs.

We design a Turing machine *M'* that on input *i* will create two Turing machines *M_a* and *M_b* as follows. The Turing machine *M_a* will ignore its input and instead just simulate *M_i* on input *i* doing whatever it does (accept if it accepts, reject if it rejects, or run forever if it runs forever). We know that it is possible to design a Turing machine that runs another machine based on what we learned about the Universal Turing Machine. The Turing machine *M_b* accepts all strings by making the start state accepting. Finally, the Turing machine *M'* runs the Turing machine *M* (for *SAME*) on input (*a, b*) and does whatever it does (accept or reject).

We now show that *M'* is a Turing machine for *D* that halts on all inputs. We prove the following three facts:

We know that *M'* halts on all inputs as the machine *M* halts on all inputs.

If $i \in D$, then *M_i* accepts *i* (by the definition of *D*) which means that *M_a* will always accept and hence $L(M_a) = \Sigma^*$. Moreover, $L(M_b) = \Sigma^*$ by design, so $L(M_a) = L(M_b)$ and thus *M* will accept, causing *M'* to accept.

If $i \notin D$, then *M_i* rejects *i* (by the definition of *D*) which means that *M_a* will always reject and hence $L(M_a) = \emptyset \neq \Sigma^* = L(M_b)$. Thus, *M* will reject, causing *M'* to reject.

Putting this together, gives us that $i \in D$ if and only if *M'* accepts, making *M'* a Turing machine for *D* that halts on all inputs. However, we already proved that *D* is not recursive, which gives us a contradiction. Thus, *SAME* is not recursive.

At the heart of this reduction is the design of the Turing machine *M_a* that ignores

its input and simulates M_i on input i . It is designed specifically so that for any $i \in D$ the Turing machine M_a will always accept giving $L(M_a) = \Sigma^*$ and for all $i \notin D$ the Turing machine M_a always rejects giving $L(M_a) = \emptyset$. This stark dichotomy allows us to use M_a in the reduction to distinguish whether a given string is in D or not by using a Turing machine for *SAME*.

More specifically, we have set up the Turing machines M_a and M_b that are given as input to the Turing machine for *SAME* as follows:

$i \in D?$	M_i accepts $i?$	$L(M_a)$	$L(M_a) = L(M_b)?$	M accepts	M' accepts
Yes	Yes	Σ^*	Yes	Yes	Yes
No	No	\emptyset	No	No	No

The proof then takes the following form:

1. Assume for a contradiction that some Turing Machine M accepts *SAME* and halts on all inputs.
2. Build a new machine M' that on input i halts and accepts if $i \in D$ and halts and rejects if $i \notin D$. It does so by running M (the Turing machine for *SAME*) on the Turing machines M_a and M_b described above. (This is the part of the proof that you will have to adapt, using a table similar to the one above, to solve most of the exercises.) Since M halts on all inputs, so must M' .
3. Since we showed in the last step that M' is a Turing machine for D that halts on all inputs, this shows that D is recursive. This is a contradiction, so *SAME* can not be recursive.

14.9 THE HALTING PROBLEM

Another classic problem that is known to be uncomputable is the **Halting Problem**. Simply put, the Halting Problem asks whether a given computation will ever halt or if it will run forever. By showing that this problem is not recursive we know, for example, that it is impossible to design a compiler that can always detect infinite loop bugs in arbitrary code.

Next, we define the Halting Problem and show that it is undecidable as well.

Definition 14.7

$$H = \{(i, x) : M_i \text{ halts on input } x\}.$$

Notice that this definition is about whether a Turing machine halts or not on a given input, making it about the behavior of the Turing machine beyond whether it accepts certain strings or not. This makes the reduction a little more complicated.

Theorem 14.7

The language H is not recursive.

Proof

We will assume that H is recursive for a contradiction. Then there must exist a Turing machine M for H that halts on all inputs. We will use M to design a Turing machine for D that halts on all inputs.

We design a Turing machine M' that on input i will compute the encoding i' of a Turing machine $M_{i'}$ that will behave the same as M_i , except if M_i is about to halt and reject $M_{i'}$ will go into an infinite loop. Next, M' will run the Turing machine M on the input (i', i) . We can imagine designing a Turing machine that automatically modifies another Turing machine to behave in the above manner since it is easy to add a state for going into an infinite loop by, for example, moving right on the tape forever. We also know that it is possible to design a Turing machine that runs another machine based on what we learned about the Universal Turing Machine.

We now show that M' is a Turing machine for D that halts on all inputs. We show the following three parts:

The Turing machine M' must halt on all inputs as the machine M halts on all inputs.

If $i \in D$, then M_i accepts i (by the definition of D), so $M_{i'}$ halts on input i , and thus it must be the case that (i', i) is accepted by M , so M' accepts.

If $i \notin D$, then M_i rejects i (by the definition of D), so by design $M_{i'}$ will reject i by running forever, and thus (i', i) is rejected by M , hence M' rejects.

The last two parts together give us that $i \in D$ if and only if M' accepts, making M' a Turing machine for D that halts on all inputs. However, we already proved that D is not recursive, which gives us a contradiction.

Thus, H is not recursive.

In practical terms:

It is impossible to design a program that can tell if another program will ever terminate.

The Halting Problem is one of the classic problems in computer science that is uncomputable.

14.10 CLASSES OF LANGUAGES

When studying the difficulty of solving a problem, computer scientists use the notion of a class of problems.

Definition 14.8

A **class** of problem is a set of languages that share a common characteristic in terms of how hard it is to compute them.

Example 14.2

Some examples of classes of problems that we have studied in this book include the regular languages (Chapter 2), the context-free languages (Chapter 12), the recursive and recursively enumerable languages (Chapter 14).

In conclusion, here is a diagram with the containment properties of these classes.




Figure 14.4: All classes of languages in this book

14.11 CHAPTER SUMMARY AND KEY CONCEPTS

- **Variations** of Turing-machine models include ones with **multiple tracks** and **multiple tapes**. These models are no more powerful than the standard Turing machine model.
- The **Church-Turing thesis** states that all mechanical computation can be performed by a Turing machine.
- The **Universal Turing machine** is one that takes as input the encoding of a Turing machine and an input and simulates that Turing machine on that input.
- A language is called **recursive** or **computable** or **decidable** if there exists a Turing machine that accepts it and halts on all inputs.

¹You have learned that there are problems unsolvable by any computer—dragons!—hence the cover of this book.

- A language is called **recursively enumerable** if there exists a Turing machine that accepts it.
- A language is recursive if and only if its complement is recursive.
- A language is recursive if and only if it and its complement are recursively enumerable.
- The **diagonal language** D and its complement can be shown to be non-recursive using **diagonalization**.
- **Reductions** are a way of showing that a language is not recursive. To do this, you must reduce from a problem known not to be recursive to the problem being shown is not recursive.
- Many problems, such as determining (1) whether a Turing machine (program) will accept a given input, (2) whether two Turing machines (programs) accept the same language, and (3) whether a Turing machine (program) will ever halt are all known to be non-recursive.
- A **class** of problem is a set of languages that have the same computational difficulty. These include **regular**, **context-free**, **recursive**, and **recursively enumerable** languages.

EXERCISES

Read the entries in the Stanford Encyclopedia of Philosophy on Computability (<https://plato.stanford.edu/entries/computability/>) and Alan Turing (<https://plato.stanford.edu/entries/turing/>) and other sources to give short answers to the following questions:

- 14.1 Who designed the Lambda Calculus?
- 14.2 What was Gödel's Incompleteness Theorem?
- 14.3 What was the *entscheidungsproblem*?
- 14.4 What was the title of the paper in which Turing introduced the Turing machine?
- 14.5 What is the Turing Test?
- 14.6 Do you think that a Turing machine can simulate a human brain?

Prove that the following problems are undecidable using reductions. For each problem, write a sentence explaining in words what the implication of the undecidability of the language is for computer programs.

- 14.7 $ALL = \{i : L(M_i) = \Sigma^*\}$
- 14.8 $INFINITE = \{i : L(M_i) \text{ is infinite}\}$